

MAE 5032 High Performance Computing: Methods and Practices

Lecture 12: Numerical Analysis Basics

Ju Liu

Department of Mechanics and Aerospace Engineering
liuj36@sustech.edu.cn



Floating-point arithmetic

Floating-point numbers

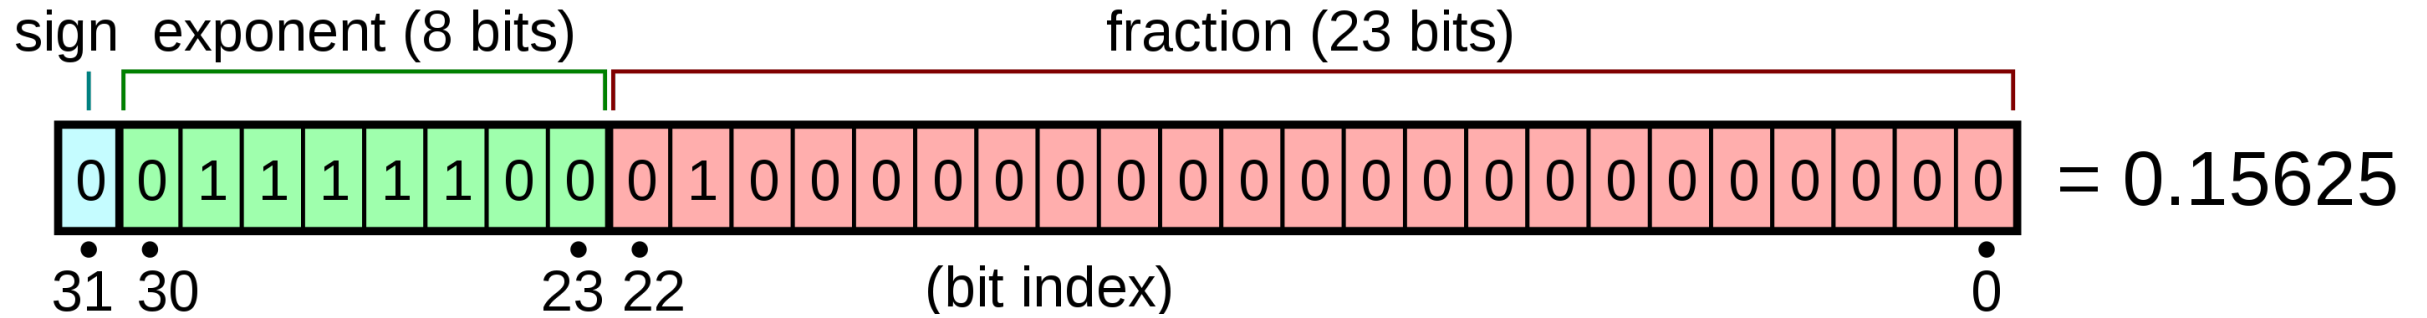
- Computers use a finite number of bits to represent numbers. Thus only a finite number of numbers can be represented.

$$x = \pm \sum_{i=0}^{t-1} d_i \beta^{-i} \beta^e = (1 \times 2^{-0} + 1 \times 2^{-1} + 0 \times 2^{-2} + \dots + 1 \times 2^{-23}) \times 2^1$$
$$\approx 1.5707964 \times 2 = 3.1415928$$

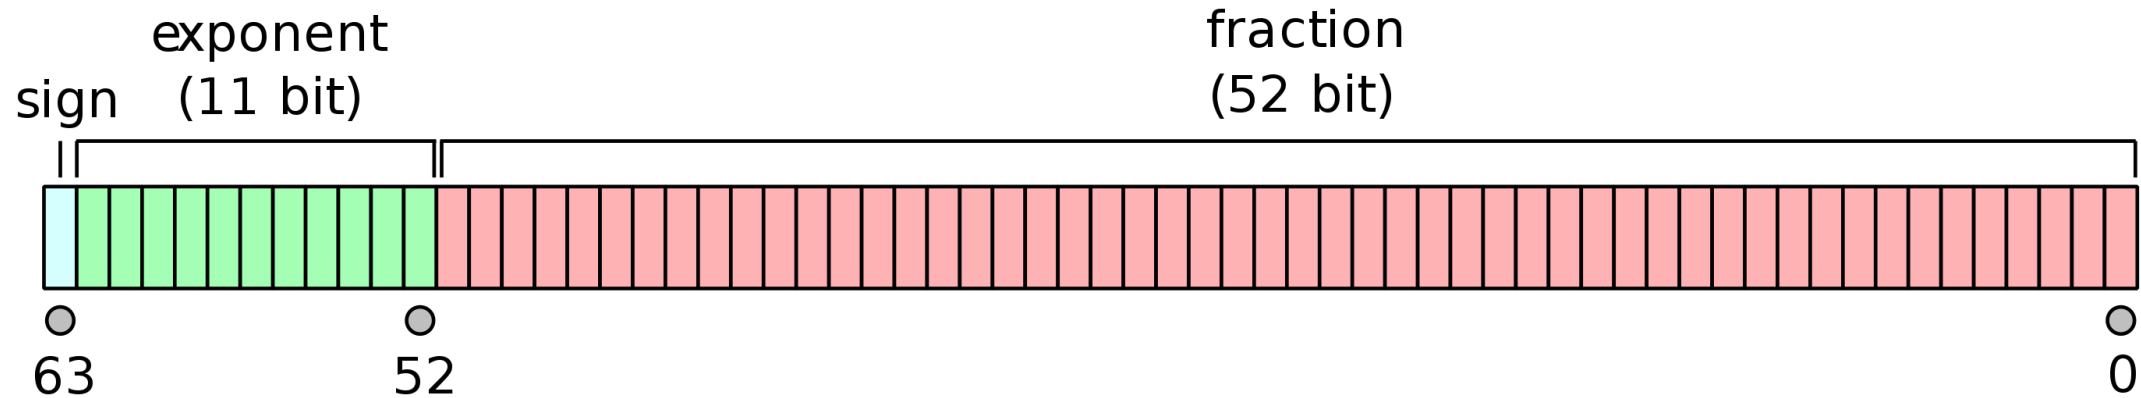
- One bit for sign (unsigned number exists)
- β is the base of the number system (2, 10, 16, etc.)
- t is the significand precision
- $0 \leq d_i \leq \beta - 1$ is the digits of the mantissa
- $L \leq e \leq U$ is the signed exponent

	β	t	L	U
IEEE single (32 bit)	2	24	-126	127
IEEE double (64 bit)	2	53	-1022	1023
Old Cray 64bit	2	48	-16383	16384
IBM mainframe 32 bit	16	6	-64	63
packed decimal	10	50	-999	999

Floating-point numbers



IEEE single 32-bit




IEEE double 64-bit

Floating-point numbers

- Normalized number require that the first digit in the mantissa to be nonzero.
- In binary system, the nonzero number is 1, thus we get one free bit for mantissa.

$$x = \pm \sum_{i=0}^t d_i \beta^{-i} \beta^e = (1 \times 2^{-0} + 1 \times 2^{-1} + 0 \times 2^{-2} + \dots 1 \times 2^{-23}) \times 2^1$$

$$\approx 1.5707964 \times 2 = 3.1415928$$


- One bit for sign (unsigned number exists)
- β is the base of the number system (2, 10, 16, etc.)
- t is the significand precision
- $0 \leq d_i \leq \beta - 1$ is the digits of the mantissa
- $L \leq e \leq U$ is the signed exponent

	β	t	L	U
IEEE single (32 bit)	2	24	-126	127
IEEE double (64 bit)	2	53	-1022	1023
Old Cray 64bit	2	48	-16383	16384
IBM mainframe 32 bit	16	6	-64	63
packed decimal	10	50	-999	999

Floating-point numbers

- Normalized number require that the first digit in the mantissa to be nonzero.
- In binary system, the nonzero number is 1, thus we get one free bit for mantissa.

$$x = \pm \sum_{i=0}^{t-1} d_i \beta^{-i} \beta^e = (1 \times 2^{-0} + 1 \times 2^{-1} + 0 \times 2^{-2} + \dots 1 \times 2^{-23}) \times 2^1$$
$$\approx 1.5707964 \times 2 = 3.1415928$$

- Underflow level β^L
- Overflow level $(1 - \beta^{-t})\beta^{U+1}$

	β	t	L	U
IEEE single (32 bit)	2	24	-126	127
IEEE double (64 bit)	2	53	-1022	1023
Old Cray 64bit	2	48	-16383	16384
IBM mainframe 32 bit	16	6	-64	63
packed decimal	10	50	-999	999

Representation error

- Error between a number x and its floating-point representation \tilde{x} :
 - absolute $|x - \tilde{x}|$
 - relative $\frac{|x - \tilde{x}|}{|x|}$
- Equivalently, sometimes we say $\tilde{x} = x(1 \pm \varepsilon)$
- IEEE 754 standard gives five rounding rules, two of them are
 - truncation: $\varepsilon = \beta^{-t}$
 - rounding: $\varepsilon = 0.5 \beta^{-t}$
- The maximum relative error is called the machine precision. (Calculate it for a 64-bit double.)

Example: decimal numerical system (i.e. $\beta=10$), $t = 3$,
 $x = 0.1256$, then $\tilde{x}_r = 0.126$ or $\tilde{x}_t = 0.125$.

Addition

- Steps for addition of floating-point numbers
 - align exponents
 - add mantissas
 - adjust exponent to normalize the result

Example: $123456.7 + 101.7654 = (1.234567 \times 10^5) + (1.017654 \times 10^2)$
 $= (1.234567 \times 10^5) + (0.001018 \times 10^5)$
 $= 1.235585 \times 10^5$

- We consider x_i and their floating-point number $\tilde{x}_i = x_i(1 + \varepsilon_i)$
To compute $s = x_1 + x_2$, the sum is represented as
$$\tilde{s} = (\tilde{x}_1 + \tilde{x}_2)(1 + \varepsilon_3) = x_1(1 + \varepsilon_1)(1 + \varepsilon_3) + x_2(1 + \varepsilon_2)(1 + \varepsilon_3)$$
$$\approx x_1(1 + \varepsilon_1 + \varepsilon_3) + x_2(1 + \varepsilon_2 + \varepsilon_3) \approx s(1 + 2\varepsilon)$$
- Conclusion: Errors are added

Subtraction and associativity

Example: $123457.1467 - 123456.659 \approx (1.234571 \times 10^5) - (1.234567 \times 10^5)$
 $= 4.000000 \times 10^{-1}$

- The actual result is 4.877000×10^{-1}
- Relative error of the subtraction is about 20%.
- In extreme cases, all significant numbers can be lost due to cancellation.

Example: 7-digits decimal floating-point number to calculate $(a+b)+c$ and $a+(b+c)$, with $a=1234.567$, $b=45.67834$, $c=0.0004$.

$a+b = 1280.24534$, rounds to 1280.245 ; $(a+b)+c = 1280.2454$ rounds to **1280.245**

$(b+c) = 45.67874$ rounds to 45.67874 ; $a+(b+c) = 1280.24574$ rounds to **1280.246**

Because of this, compilers will not automatically reorder the operations for FP.

You need to enable “fast-math” options.

A toy problem

Evaluate $\sum_{n=1}^{10000} \frac{1}{n^2}$. The precise value is 1.644834 in a decimal numerical system (i.e. $\beta=10$) with $t = 7$.

First term is 1.000000, so the partial sum will be greater than 1. So, for the terms that $\frac{1}{n^2} < 10^{-6}$, their contribution to the sum will be ignored.

Floating point sum is 1.644725: 4 correct digits.

Solution: sum in reverse order.

Unstable algorithm

Consider the recurrence

$$y_n = \int_0^1 \frac{x^n}{x+5} dx. \text{ We may deduce that } y_n = \frac{1}{n} - 5y_{n-1}.$$

The initial value of the recurrence is $y_0 = \log(6) - \log(5) = 1.82|322 \times 10^{-1}$.

Consider a decimal numerical system (i.e. $\beta=10$) with $t = 3$.

FP arithmetic

$$\tilde{y}_0 = 1.82 \times 10^{-1}$$

$$\tilde{y}_1 = 9.00 \times 10^{-2}$$

$$\tilde{y}_2 = 5.00 \times 10^{-2}$$

$$\tilde{y}_3 = 8.30 \times 10^{-2}$$

$$\tilde{y}_4 = -1.65 \times 10^{-1}$$

Correct value

$$y_0 = 1.82 \times 10^{-1}$$

$$y_1 = 8.84 \times 10^{-2}$$

$$y_2 = 5.80 \times 10^{-2}$$

$$y_3 = 4.31 \times 10^{-2}$$

$$y_4 = 3.43 \times 10^{-2}$$

Let $\tilde{y}_n = y_n + \varepsilon_n$. Then $\tilde{y}_n = \frac{1}{n} - 5 \tilde{y}_{n-1} = \frac{1}{n} - 5y_{n-1} + 5\varepsilon_{n-1} = y_n + 5\varepsilon_{n-1}$.

$\varepsilon_n = 5\varepsilon_{n-1}$. Error grows exponentially.

Summary

- Arithmetic on computer is done based on floating-point numbers, and thus simple calculations like addition may lead to error.
- Mathematically equivalent operations may not remain equivalent on computers, which could affect parallel computations.
- Algorithms need to be carefully designed to control the floating-point error
- Stability analysis is needed for ALL operations on computers.

Reference: Introudction to High Performance Scientific Computing by Victor Eijkhout, Chapter 2.

Linear Algebra

A known unusable algorithm

To solve $Ax = b$, one may think of the Cramer's rule:

$$x_i = \frac{\begin{vmatrix} a_{11} & a_{12} & \dots & a_{1i-1} & b_1 & a_{1i+1} & \dots & a_{1n} \\ a_{21} & & \dots & & b_2 & & \dots & a_{2n} \\ \vdots & & & & \vdots & & & \vdots \\ a_{n1} & & \dots & & b_n & & \dots & a_{nn} \end{vmatrix}}{|A|}$$

The time complexity is $O(n!)$

Recall that on a single CPU of TaiYi, we can achieve around 76.8 Gflops.

This algorithm is too expensive! Solving a 20-by-20 matrix problem may take 1 year!

Direct method

To solve $Ax = b$, one may do the old Gaussian elimination method:

$$\begin{pmatrix} 6 & -2 & 2 \\ 12 & -8 & 6 \\ 3 & -13 & 3 \end{pmatrix} x = \begin{pmatrix} 16 \\ 26 \\ -19 \end{pmatrix}$$

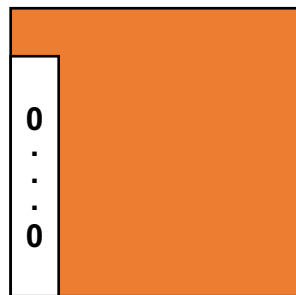
$$\left[\begin{array}{ccc|c} 6 & -2 & 2 & 16 \\ 12 & -8 & 6 & 26 \\ 3 & -13 & 3 & -19 \end{array} \right] \longrightarrow \left[\begin{array}{ccc|c} 6 & -2 & 2 & 16 \\ 0 & -4 & 2 & -6 \\ 0 & -12 & 2 & -27 \end{array} \right] \longrightarrow \left[\begin{array}{ccc|c} 6 & -2 & 2 & 16 \\ 0 & -4 & 2 & -6 \\ 0 & 0 & -4 & -9 \end{array} \right]$$

Algorithmic complexity is $O(n^3)$.

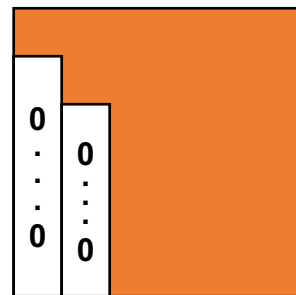
Gaussian Elimination algorithm

- Add multiples of each row to later rows to make A upper triangular
- Solve resulting triangular system $Ux = c$ by substitution

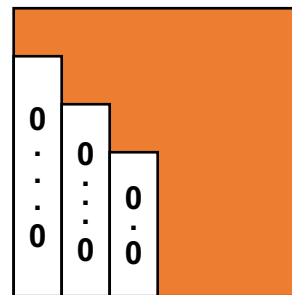
```
... for each column i
... zero it out below the diagonal by adding multiples of row i to later rows
for i = 1 to n-1
  ... for each row j below row i
  for j = i+1 to n
    ... add a multiple of row i to row j
    tmp = A(j,i);
    for k = i to n
      A(j,k) = A(j,k) - (tmp/A(i,i)) * A(i,k)
```



After i=1

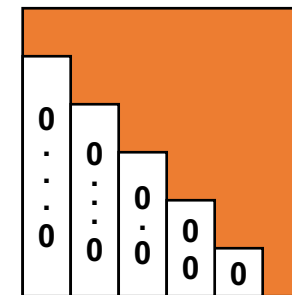


After i=2



After i=3

...



After i=n-1

Gaussian Elimination algorithm: math kernel

- LU factorization: if the above algorithm completes, we get $A = LU$. $\square = \triangle * \nabla$
- Thus $Ax=b$ becomes $LUx=b$.
 - We solve $Ly=b$ first and solve $Ux=y$ next.
- Solving $Ax=b$ using GE:
 - Factorize $A = L*U$ using GE (cost = $\frac{2}{3} n^3$ flops)
 - Solve $L*y = b$ for y , using substitution (cost = n^2 flops)
 - Solve $U*x = y$ for x , using substitution (cost = n^2 flops)

Roundoff control

Consider a system

$$\begin{bmatrix} \varepsilon & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 + \varepsilon \\ 2 \end{bmatrix}.$$

The exact solution is

$$\begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}.$$

Perform Gaussian elimination

$$\begin{bmatrix} \varepsilon & 1 \\ 0 & 1 - \frac{1}{\varepsilon} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 + \varepsilon \\ 1 - \frac{1}{\varepsilon} \end{bmatrix}.$$

We can do a “back-substitution” by solving x_2 first and solve x_1 next:

$$x_2 = 1 \Rightarrow x_1 = 1.$$

Roundoff control (cont.)

Suppose ε is smaller than the machine precision. $1 - \frac{1}{\varepsilon}$ becomes $-\frac{1}{\varepsilon}$, and $1 + \varepsilon$ becomes 1.

The Gaussian elimination

$$\begin{bmatrix} \varepsilon & 1 \\ 0 & 1 - \frac{1}{\varepsilon} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 + \varepsilon \\ 1 - \frac{1}{\varepsilon} \end{bmatrix}$$

becomes

$$\begin{bmatrix} \varepsilon & 1 \\ 0 & -\frac{1}{\varepsilon} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 \\ -\frac{1}{\varepsilon} \end{bmatrix}$$

We can do a “back-substitution” by solving x_2 first and solve x_1 next:

$$x_2 = 1 \Rightarrow x_1 = 0 .$$

Machine round-off error has a dramatic impact on the results. We call this **numerical instability**.

Pivoting in Gaussian Elimination

- During the LU factorization, when the diagonal is small, swap the row (or column) to avoid division by small numbers.
 - This is referred to as pivoting in GE
 - We choose the largest possible pivot.
 - In fact, we always choose the largest possible value as the pivot in GE.
 - There is always a nonzero pivot if the matrix is non-singular.

Roundoff control with pivoting

Consider the system again,

$$\begin{bmatrix} \varepsilon & 1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 + \varepsilon \\ 2 \end{bmatrix}.$$

Pivot the row

$$\begin{bmatrix} 1 & 1 \\ \varepsilon & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 2 \\ 1 + \varepsilon \end{bmatrix}.$$

Perform Gaussian elimination

$$\begin{bmatrix} 1 & 1 \\ 0 & 1 - \varepsilon \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 2 \\ 1 - \varepsilon \end{bmatrix}.$$

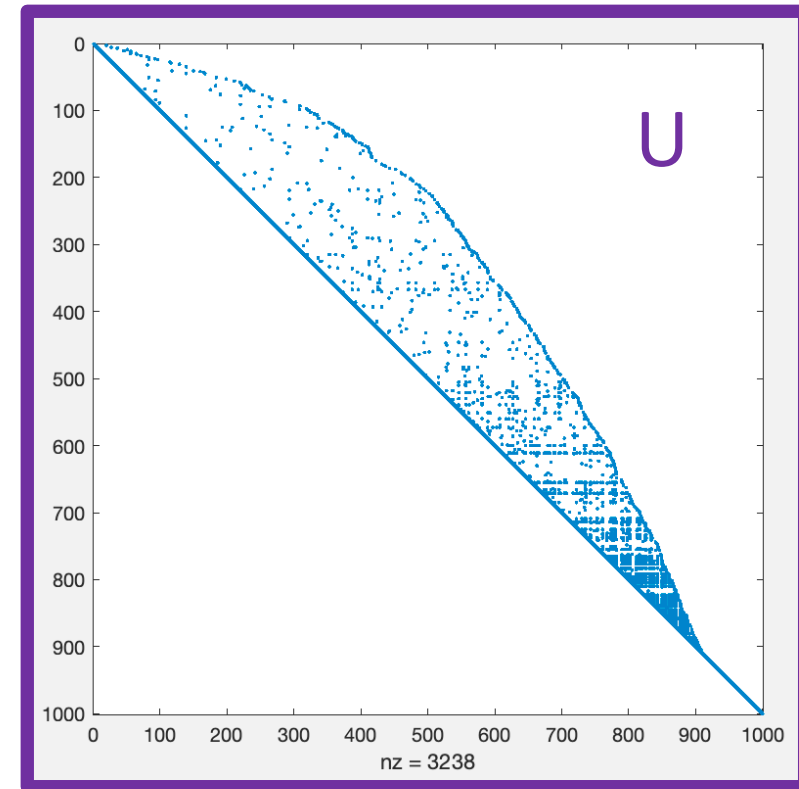
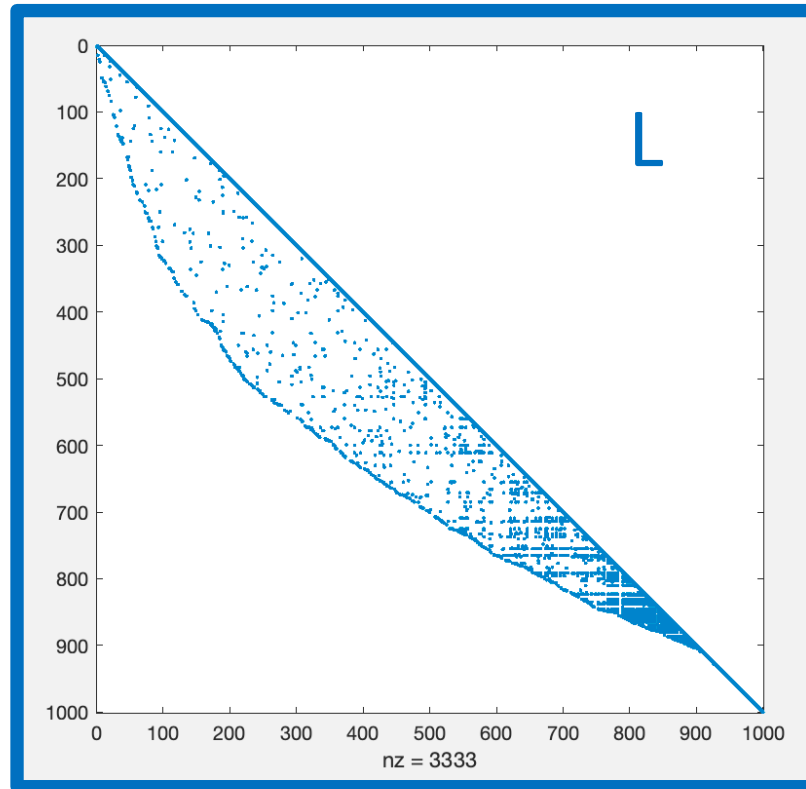
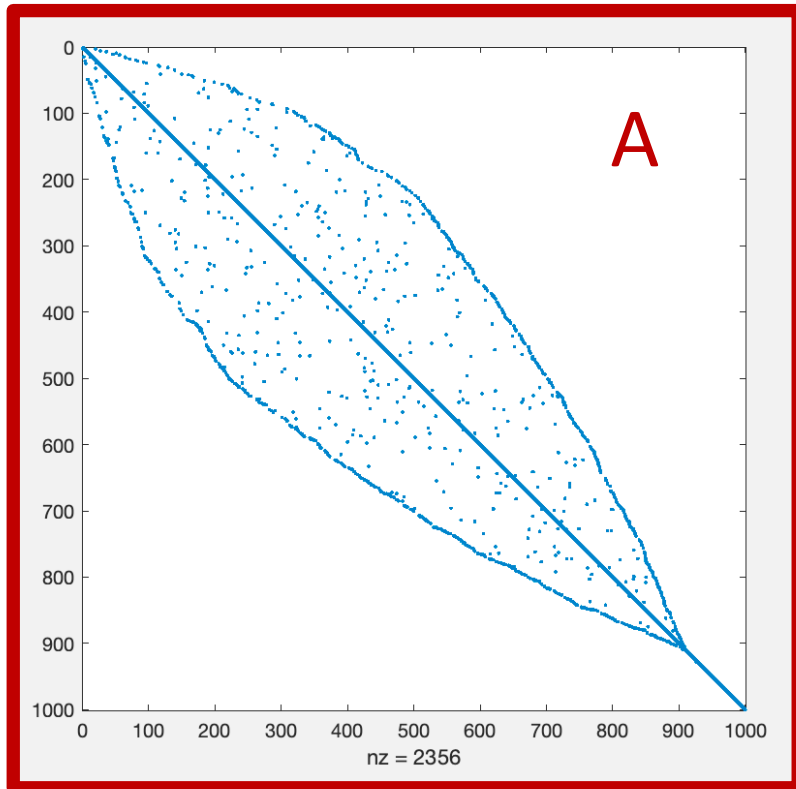
If ε is small,

$$\begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 2 \\ 1 \end{bmatrix}.$$

Back substitution: $x_2 = 1 \Rightarrow x_1 = 1$.

The fill-in phenomenon

- LU factorization of a sparse matrix typically leads to additional nonzeros. This phenomenon is called **fill-in**.
- The memory requirement can thus become a bottleneck if one wants to use direct method to solve a sparse matrix.



1990 Nobel Prize in Economics



Our models strained the computer capabilities of the day [1950s]. I observed that most of the coefficients in our matrices were zero; i.e., the nonzeros were sparse in the matrix.

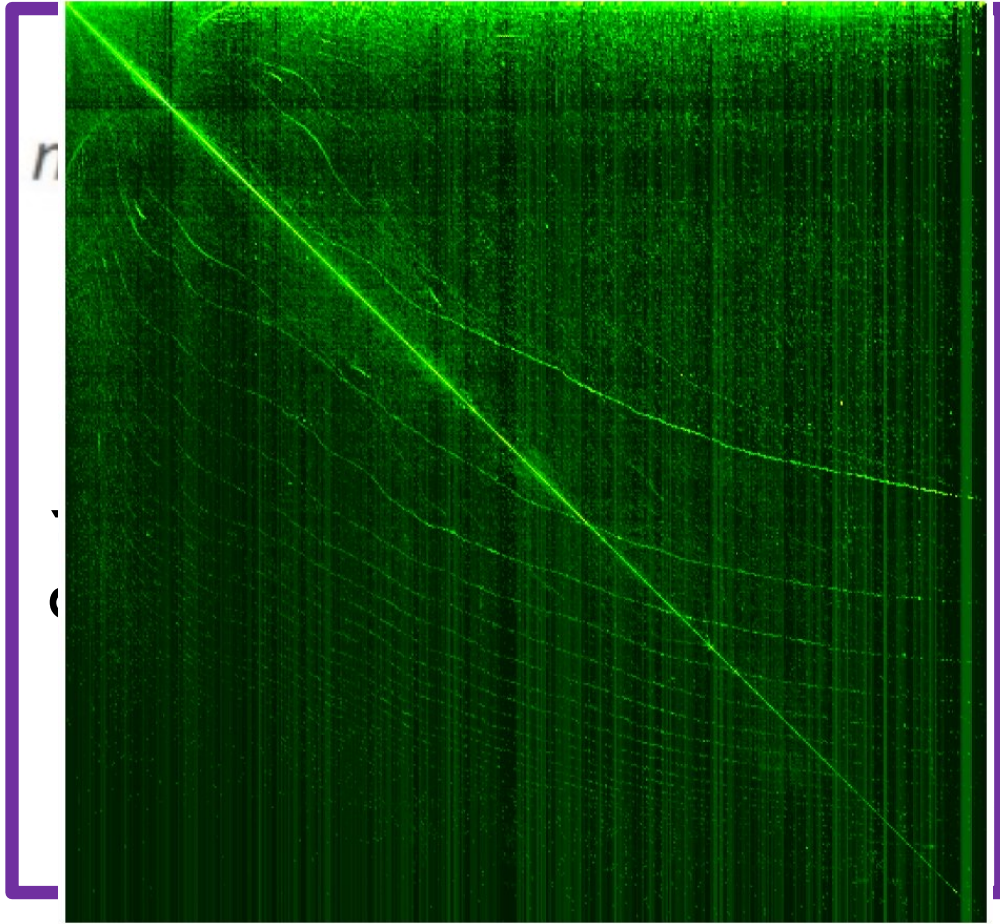
-- Harry Markowitz

Sparse matrices arise in many applications:

- simulating physics
- analyzing images
- web page ranking in search engines
-

Examples

Products

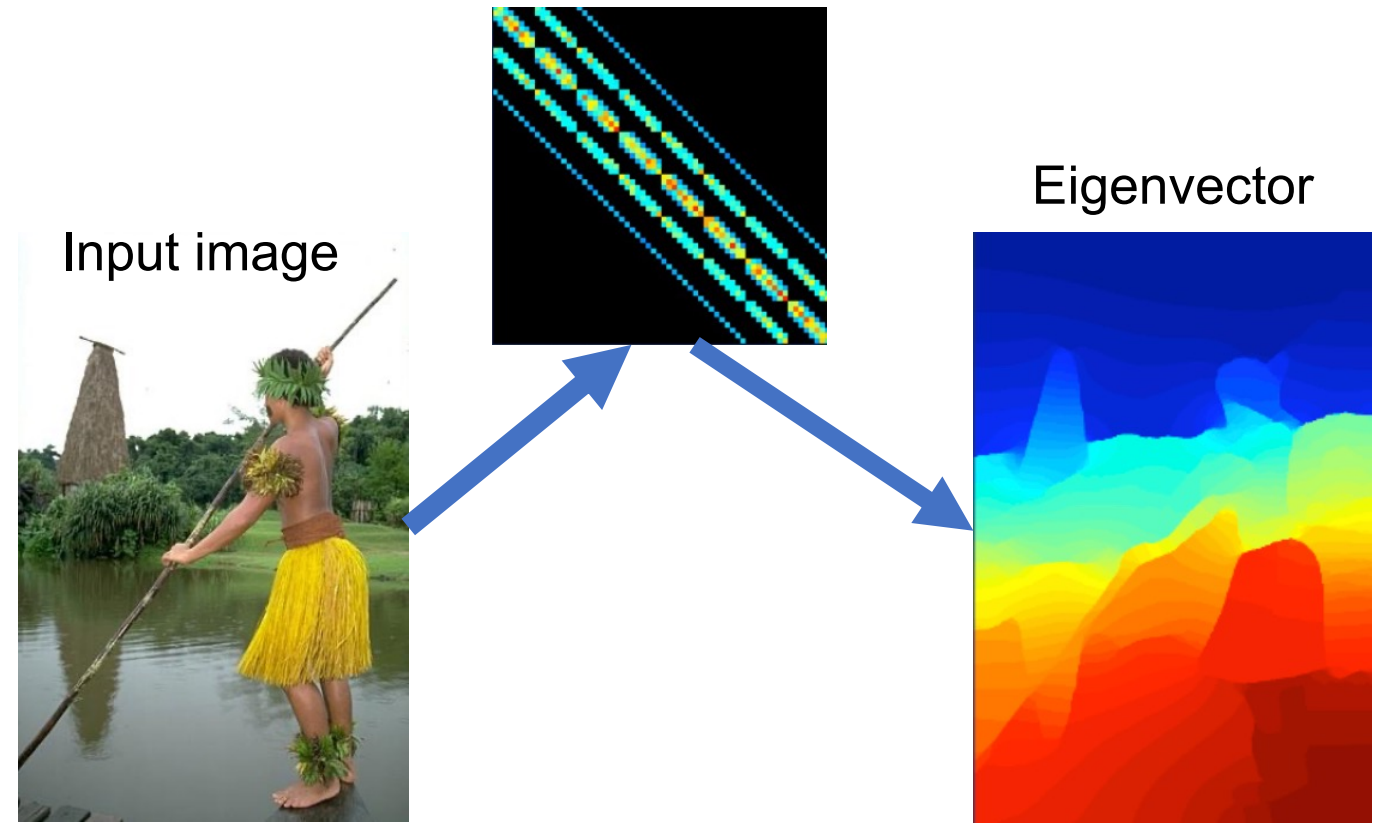


Recommendation Matrix

Image segmentation – identify the object boundaries in an image.

Find the eigenvalue of the affinity matrix.

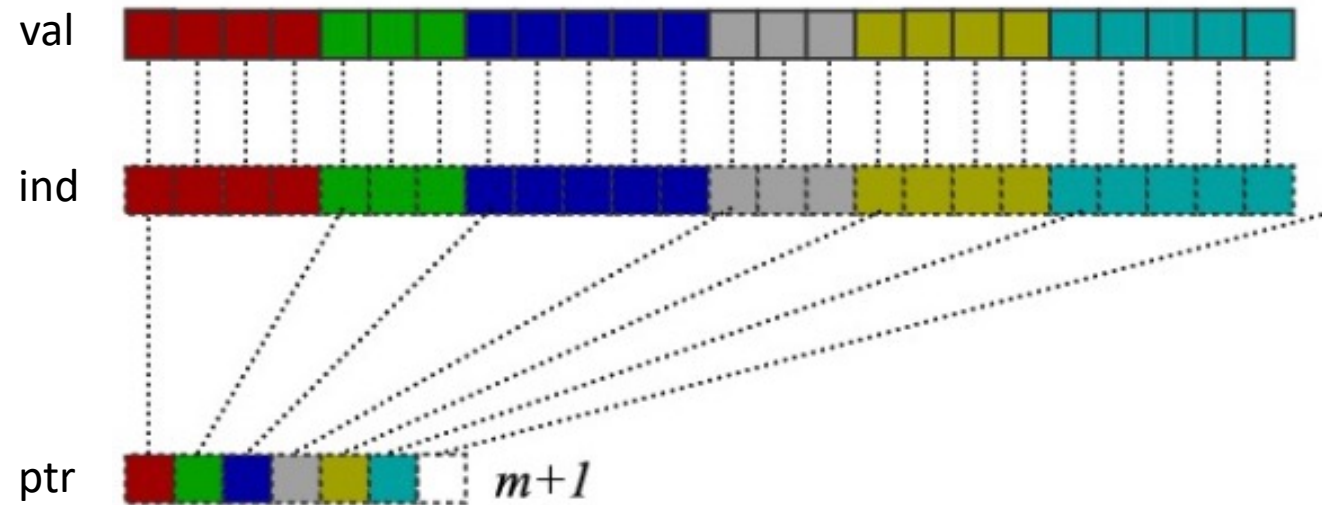
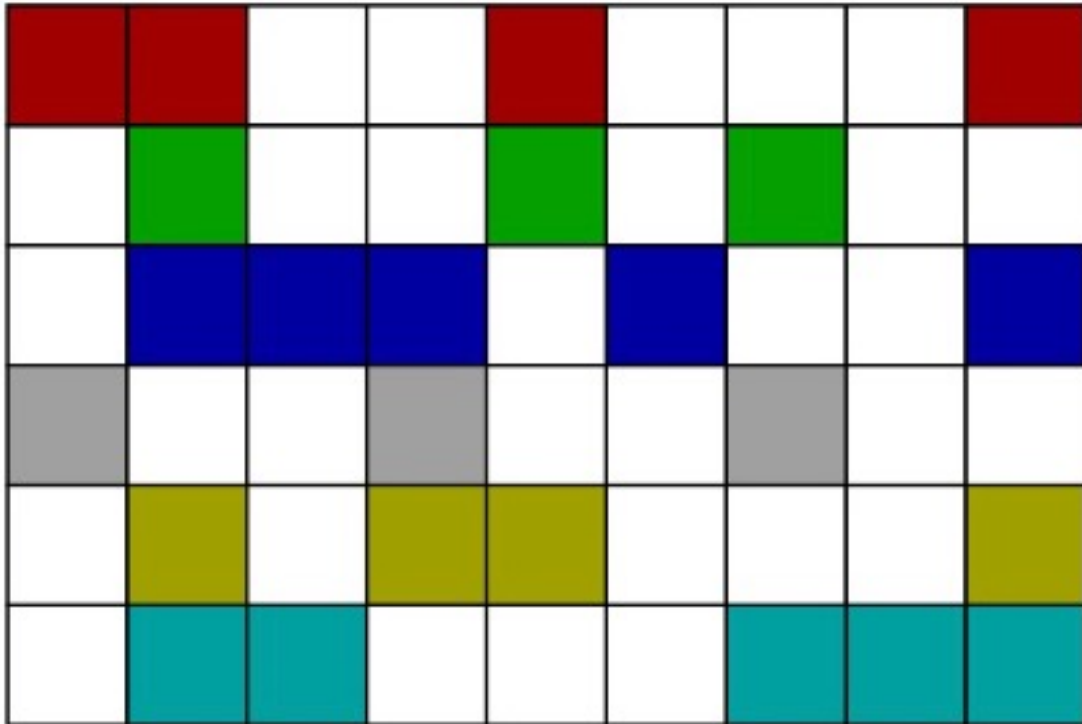
Efficient, High-quality image contour detection by Catanzaro, Su, et al. 2009.



Data structure

- Use a proper data structure to represent the matrix by saving the nonzeros only.
- Compressed Row Storage (CRS) is the most widely used format.

Example:



Data structure

- Use a proper data structure to represent the matrix by saving the nonzeros only.
- Compressed Row Storage (CRS) is the most widely used format.

Example:

$$A = \begin{pmatrix} 10 & 0 & 0 & 0 & -2 & 0 \\ 3 & 9 & 0 & 0 & 0 & 3 \\ 0 & 7 & 8 & 7 & 0 & 0 \\ 3 & 0 & 8 & 7 & 5 & 0 \\ 0 & 8 & 0 & 9 & 9 & 13 \\ 0 & 4 & 0 & 0 & 2 & -1 \end{pmatrix}$$

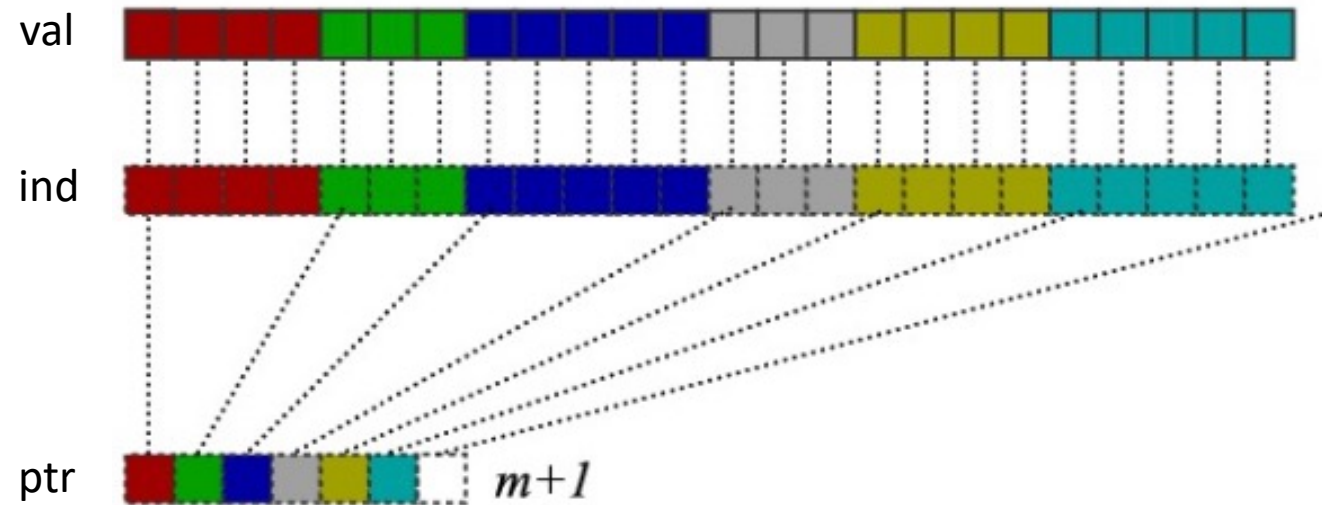
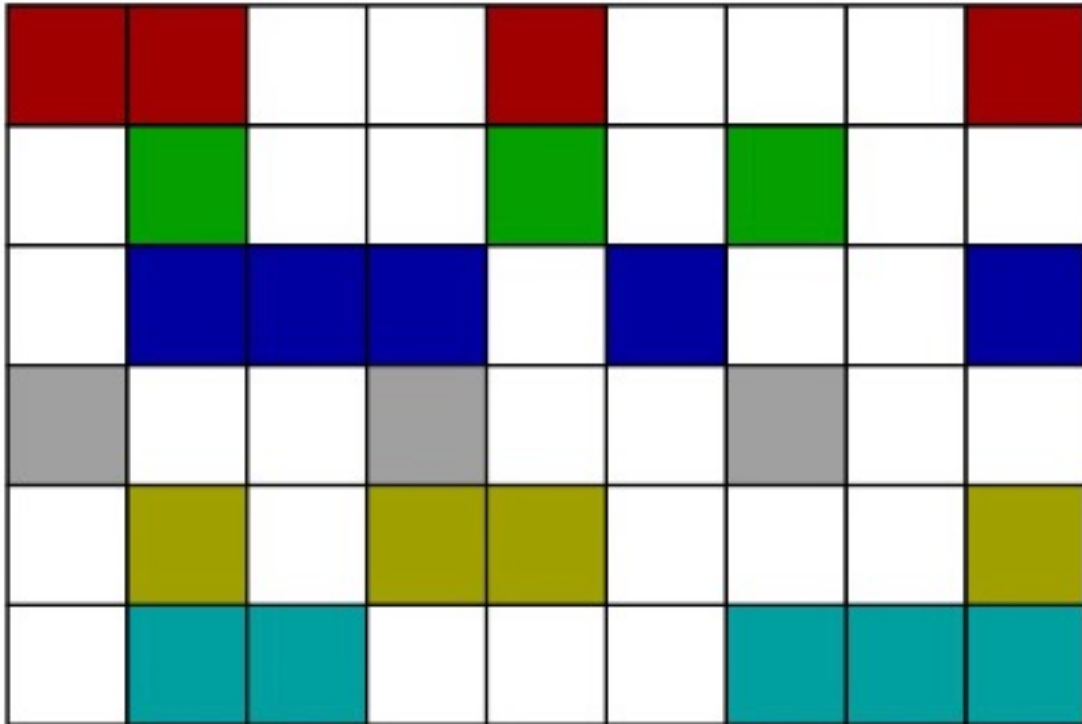
val	10	-2	3	9	3	7	8	7	3 ... 9	13	4	2	-1
ind	0	4	0	1	5	1	2	3	0 ... 4	5	1	4	5
ptr	0	2	5	8	12	16	19						

Preallocation of the memory space can be critical! Otherwise, the initial use of the matrix can be extremely slow.

Data structure

- Use a proper data structure to represent the matrix by saving the nonzeros only.
- Compressed Row Storage (CRS) is the most widely used format.

Example:



```
y(i) ← y(i) + A(i,j) x(j)
for each row i
  for k = ptr[i] to ptr[i+1]
    y[i] += val[k] * x[ind[k]]
```

Stationary iterative method

- To solve $Ax = b$ with the matrix A being a sparse matrix, we want to devise an iterative method

$$Mx_{k+1} = Nx_k + f.$$

such that the solution of $Ax = b$ also satisfies $Mx = Nx + f$.

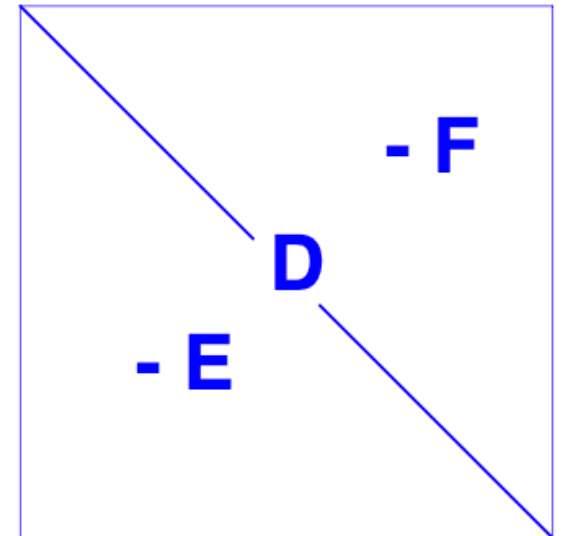
- A lot iterations can be devised based on the decomposition $A = D - E - F$.

- Jacobi iteration: $Dx_{k+1} = (E + F)x_k + b$.

- Gauss-Seidel iteration: $(D - E)x_{k+1} = Fx_k + b$.

- Successive Over-Relaxation (SOR):

$$(D - \omega E)x_{k+1} = (\omega F + (1 - \omega)D)x_k + \omega b$$



Stationary iterative method

- If the iteration converges, it converges to the linear system solution.
- In one iteration, in general, we only need to perform (1) matrix-vector multiplication; (2) perhaps solving a simpler matrix problem (usually can be achieved component-wisely). The cost of one iteration is $O(n^2)$.
- Stopping criterion:
 - monitor the residual $\|b - Ax_k\|$
 - monitor the relative change $\|x_{k+1} - x_k\|$
- It is often the case that we can achieve desirable accuracy in less than n iteration, meaning the iterative method is rather competitive, when compared against direct method.
- Is the method guaranteed to converge? How fast does it converge?
There are theories to support this. Often is problem-specific.

Krylov iterative method

- Aleksey N. Krylov (1863-1945) is a Russian naval engineer and applied mathematician.
- He published around 300 papers, topics include shipbuilding, magnetism, artillery, mathematics, astronomy, etc.
- He built the first machine in Russia for integrating ODEs.
- In 1931, he published a paper on what is now called the **Krylov subspace** and **Krylov subspace methods**.



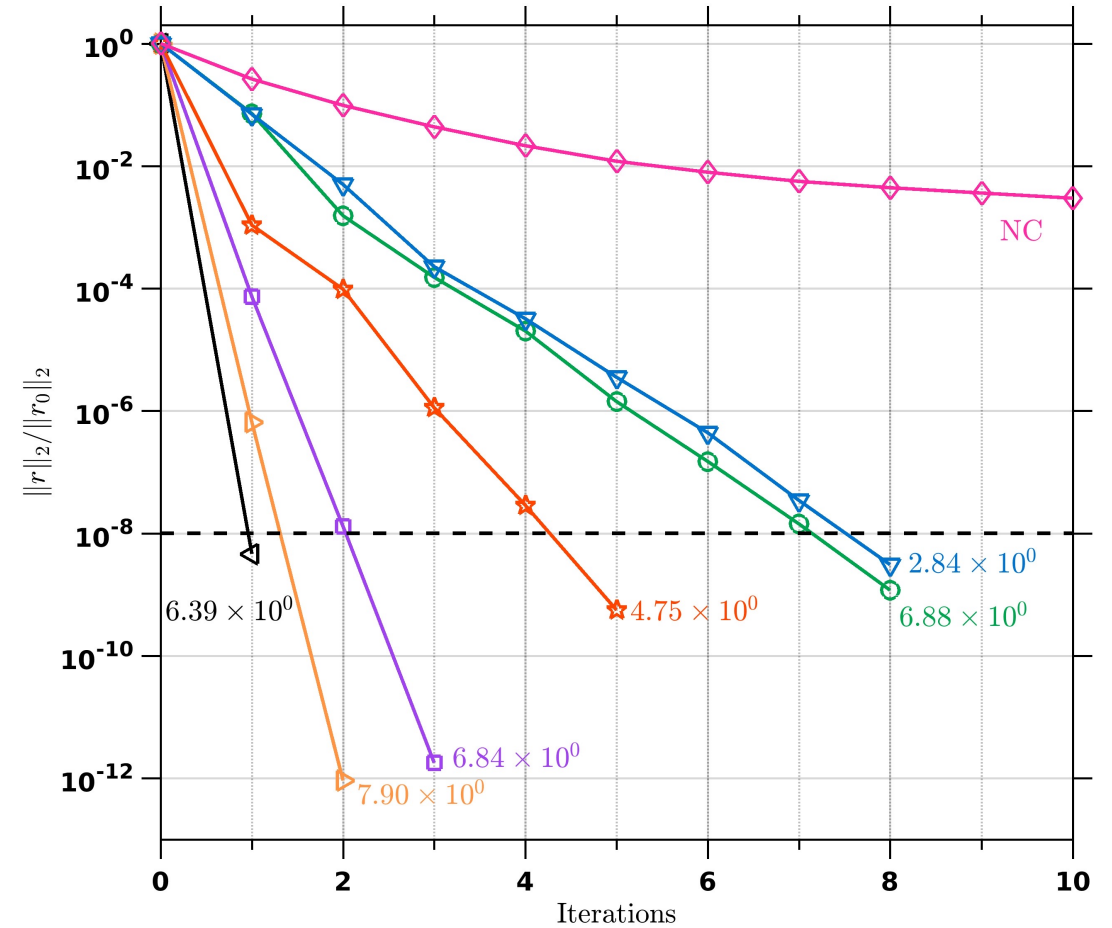
Krylov iterative method

- We are considering solving $Ax = b$, in which A is large, sparse, and possibly with irregular structures.
- The idea is to solve a least-square problem: $\min_{z \in K} \|b - Az\|$
- Krylov subspace: $K_m(A, v_1) = \text{span}\{v_1, Av_1, \dots, A^{m-1}v_1\}$
- There are very fast algorithms that can locate the vector that minimizes $\|b - Az\|$, which are referred to as the Krylov method
 - Use Arnoldi to orthogonalize the subspace
 - Use QR-decomposition to find the minimizer (NOT Least Square Method!!!)
 - Include GMRES, Conjugate Gradient (CG), BICGSTAB, MINRES, etc.
 - All have been implemented in free, open-source libraries, do not do it yourself.

Krylov iterative method

- We can modify the linear system by considering solving
$$P_l A x = P_l b,$$
or $AP_r y = b$, with $x = P_r y$ or $P_l AP_r x = P_l b$, with $x = P_r y$
- The matrices P_r and P_l are left and right preconditioners, which may accelerate the Krylov iteration.

Examples: Jacobi, ILU, etc.



Summary

- Linear algebra problem is ubiquitous.
- The condition number measures the property of a matrix.
- Direct method such as LU factorization with pivoting can solve a problem in a very robust way.
 - Its computational cost is $O(n^3)$.
 - For sparse matrices, the factorization will lead to a fill-in phenomenon, which could cost a lot memory space.
- Iterative method strives to solve the linear problem by matrix-vector multiplications iteratively.
 - The cost per iteration is cheap $O(n^2)$ at most and $O(n)$ for very sparse matrices.
 - Iteration drives the error down to certain prescribed tolerance.
 - Preconditioner may accelerate the convergence.
 - GMRES and CG are most commonly used method for solving linear systems.

Summary

Reference: Numerical Linear Algebra by L.N. Trefethen and D. Bau, III.

Differential equations

Initial value problem

- Consider the problem

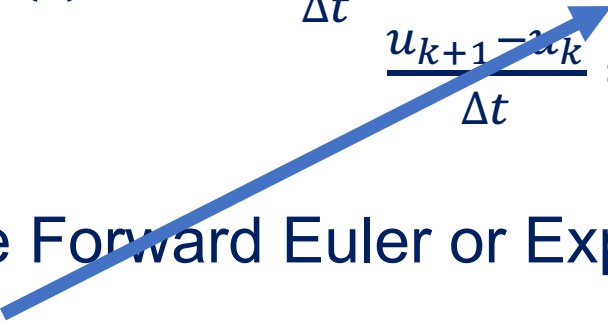
$$\frac{du}{dt}(t) = f(t, u(t)), \quad t \in (0, T), \quad \text{with } u(0) = u_0.$$

- We may turn the continuous problem into a discrete one by looking at finite time steps: $t_0 = 0, t_1, t_2, \dots, t_N = T$. Here we assume the time step size is uniform: Δt .

- Taylor series

$$u(t + \Delta t) = u(t) + u'(t)\Delta t + u''(t)\frac{\Delta t^2}{2} + \dots$$

- Use the following: $u'(t) = \frac{u(t+\Delta t) - u(t)}{\Delta t} + O(\Delta t)$, we get a discrete problem:

$$\frac{u_{k+1} - u_k}{\Delta t} = f(t_k, u_k)$$


- This is known as the Forward Euler or Explicit Euler method.
- First-order accurate.

Stability of the explicit Euler method

- Consider the problem

$$\frac{du}{dt}(t) = f(t, u(t)), \quad t \in (0, T), \quad \text{with } u(0) = u_0.$$

Let $f = -\lambda u$. The exact solution is $u(t) = u_0 e^{-\lambda t}$. It is monotonically decreasing for positive λ .

$$u_{k+1} = u_k + \Delta t f(t_k, u_k) = (1 - \lambda \Delta t) u_k = \dots = (1 - \lambda \Delta t)^{k+1} u_0$$

To have the discrete solution mimic its continuous counterpart,

$$|1 - \lambda \Delta t| < 1$$

which leads to $\Delta t < 2/\lambda$.

We call a method that has constraint on the selection of grid size to be **conditionally stable**.

Implicit Euler Method

- Consider the problem

$$\frac{du}{dt}(t) = f(t, u(t)), \quad t \in (0, T), \quad \text{with } u(0) = u_0.$$

- Taylor series

$$u(t - \Delta t) = u(t) - u'(t)\Delta t + u''(t)\frac{\Delta t^2}{2} + \dots$$

- Use the following: $u'(t) = \frac{u(t) - u(t - \Delta t)}{\Delta t} + O(\Delta t)$

We get a discrete problem:

$$\frac{u_{k+1} - u_k}{\Delta t} = f(t_{k+1}, u_{k+1})$$

- This is known as the Backward Euler or Implicit Euler method.
- First-order accurate.

Stability of the implicit Euler method

- Consider the problem

$$\frac{du}{dt}(t) = f(t, u(t)), \quad t \in (0, T), \quad \text{with } u(0) = u_0.$$

Let $f = -\lambda u$. The exact solution is $u(t) = u_0 e^{-\lambda t}$. It is monotonically decreasing for positive λ .

$$u_{k+1} = u_k + \Delta t f(t_{k+1}, u_{k+1}) \Rightarrow u_{k+1} = \frac{1}{1+\lambda\Delta t} u_k = \cdots = \left(\frac{1}{1+\lambda\Delta t}\right)^{k+1} u_0$$

To have the discrete solution mimic its continuous counterpart, there is no limit on the choice of time step size.

We call a method that has **NO** constraint on the selection of grid size to be **unconditionally stable**.

Explicit vs Implicit

- Consider the problem

$$\frac{du}{dt}(t) = f(t, u(t)), \quad t \in (0, T), \quad \text{with } u(0) = u_0.$$

Let $f = u^2$.

Forward Euler: $u_{k+1} = u_k + \Delta t f(t_k, u_k) = u_k + u_k^2$.

Backward Euler: $u_{k+1} = u_k + \Delta t f(t_k, u_{k+1}) = u_k + u_{k+1}^2$

In **implicit** method, we need to solve equations to determine the solution.

There are other options: Trapezoidal/Crank-Nicholson rule, Runge-Kutta rule, etc.

Boundary value problem

- Consider the problem

$$\frac{d^2u}{dx^2}(x) = f\left(x, u, \frac{du}{dx}\right), \quad x \in (a, b), \text{ with } u(a) = u_a \text{ and } u(b) = u_b.$$

- Taylor series

$$u(x + \Delta x) = u(x) + u'(x)\Delta x + u''(x)\frac{\Delta x^2}{2} + u'''(x)\frac{\Delta x^3}{6} + \dots$$

$$u(x - \Delta x) = u(x) - u'(x)\Delta x + u''(x)\frac{\Delta x^2}{2} - u'''(x)\frac{\Delta x^3}{6} + \dots$$

- We can get: $u''(x) = \frac{u(x+\Delta x) - 2u(x) + u(x-\Delta x)}{\Delta x^2} + O(\Delta x^2)$

We get a discrete problem:

$$\frac{u_{k+1} - 2u_k + u_{k-1}}{\Delta x^2} = f(x_k, u_k, u'_k)$$

Boundary value problem

- We get a discrete problem:

$$\frac{u_{k+1} - 2u_k + u_{k-1}}{\Delta x^2} = f(x_k, u_k, u'_k)$$

Written as a matrix problem:

$$\begin{bmatrix} 2 & -1 & \cdots & 0 \\ -1 & 2 & & 0 \\ & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 2 \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ \vdots \\ \vdots \end{bmatrix} = -\Delta x^2 \begin{bmatrix} f_1 + u_a \\ f_2 \\ \vdots \\ \vdots \end{bmatrix}$$

Matrix properties: Very sparse (tri-diagonal), symmetric, positive definite.

Summary

- We can address the differential equations by replacing the **differential** operators by their **difference** counterpart.
- There are explicit and implicit method when marching in time.
- Typically, explicit method is conditionally stable, which poses a constraint on the choice of time step size.
- Implicit method can be either conditionally stable or unconditionally stable.
- Unconditionally stable algorithms is useful for delivering steady-state solutions.
- Accuracy analysis can be achieved by standard calculus techniques, such as the Taylor expansion.
- Accuracy can be verified by numerical results if you know the exact solution.

Summary

- Manufactured solution can be an effective way in designing code verifications.
- One may plot the error in a log-log plot against the mesh size, and the error shall be in a straight line with the slope being the accuracy.

Reference: 《微分方程数值解法》 戴嘉尊，邱建贤。